


Saliency Maps & Input-Space Interpretability

LUCA DE GABRIELE

77304L

ARI5118 – DEEP LEARNING FOR COMPUTER VISION

WHY INTERPRETABILITY MATTERS




What Influenced this Decision?

High accuracy does not guarantee that reasoning is trustworthy.





THE BIOLOGICAL BLUEPRINT FOR MACHINE VISION

Primates interpret complex scenes in real-time despite slow neuronal hardware. They achieve this via a serial strategy: a 'focus of attention' scans the scene.



CONSTRUCTING SALIENCE: THE 1998 PIPELINE

How raw pixels are filtered into a single point of focus.



SALIENCY MAP




WINNER-TAKE-ALL



ATTENDED LOCATION

WHAT ARE INPUT GRADIENTS?



How does changing one pixel affect the final confidence score?

Large Change -> Important

Small Change -> Less Important

WHAT IS A SALIENCY MAP?

- A **saliency map** is the result of gathering all of the tiny “nudges” and displaying them as a single image.
- This map works in a similar way as a heat-sensing camera for the AI’s brain.
- When the AI looks at pictures of animals, the saliency maps would highlight important areas of the animals such as ears and the beak of the flamingo.
- A saliency map can turn the mathematical importance of each pixel into a visual guide that humans can understand.



VANILLA SALIENCY MAPS

- Find the **derivative of the class score** with respect to *the input image*

$$M_c(\mathbf{x}) = \partial S_c(\mathbf{x}) / \partial(\mathbf{x})$$

- Equation represents how much a tiny change in each pixel of the input $\mathbf{x}$ affects the final classification score.

- **Algorithm**

- A **forward pass** through the network is performed with the **input $\mathbf{x}$** to be able to calculate the **activation function S_c** for the **target class c** .
- A **backpropagation pass** is performed to be able to **find the gradient of S_c** with respect to the input pixels.
- Optionally, the **absolute value of the gradients** can be taken or a **limit** can be put on any outlying values to create a more visually coherent heatmap.


VANILLA SALIENCY MAPS

Original Image



Gradient
Calculation

Vanilla Saliency Map



GUIDED BACKPROPAGATION


- This is designed as an **improvement for Vanilla Saliency Map**.
 - To prioritise positive contributions to the final prediction
 - Done by modifying the standard backpropagation algorithm
- Guided backpropagation adjusts the calculation of the input gradients at the **Rectified Linear Unit (ReLU) layers**.
- *Algorithm*
 - A **standard forward pass** is performed to be able to compute activations.
 - During backpropagation, at each ReLU layer, any **negative gradient values are discarded**.
 - Effectively, the zero-masks gradients where the gradient is negative or where the input activation was negative, focusing the visualization on features that positively triggered high-level units.

GUIDED BACKPROPAGATION


Original Image



Vanilla Saliency Map



Guided Backprop



SMOOTHGRAD

- Used to “sharpen” any noisy gradient-based maps
 - Done by averaging the gradients of many versions of the same image perturbed with noise

$$\hat{M}_c(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n M_c(\mathbf{x} + \mathbf{N}(\mathbf{o}, \sigma^2))$$

n = the number of noise trials,

$\mathbf{N}(\mathbf{o}, \sigma^2)$ = the Gaussian noise with standard deviation

$\sigma \hat{M}_c$ = the smoothed gradient of $M_c(\mathbf{x})$.

$\hat{M}_c$ = the average of the model’s sensitivity measured according to slightly distorted maps of n noisy samples

- **Algorithm:**

- A sample size n and a noise level σ are selected.
- $\mathbf{N}$ perturbed images are then generated by adding random Gaussian noise to the original input $\mathbf{x}$.
- The vanilla gradient (sensitivity map) for each of the n perturbed images is calculated.
- The average of these n gradients is then calculated to produce the final sharpened saliency map.


SMOOTHGRAD

Original Image




Vanilla Saliency
+ Average

SmoothGrad Saliency




SMOOTHGRAD




SMOOTHGRAD


Noise level = 0.0



Noise level = 0.3



Noise level = 0.5



INTERACTIVE SIMULATOR

WHEN SALIENCY MAPS FAIL

- Some saliency methods produce **visualisations** that are **independent of both the model's weights and the data generation process**.


This problem is known as the **Edge Detector Problem**.

- Many saliency methods are unreliable because they **lack input invariance**.
 - Input invariance is a requirement that an attribution method needs to replicate the sensitivity of the model regardless of how the input is transformed.
- Many existing methods fail to satisfy sensitivity and implementation invariance.
- **Raw gradient-based maps** are **visually noisy, highlighting pixels** that can appear randomly to the human eye.

WHEN NOT TO USE SALIENCY MAPS


SCENARIO	LIMITATION	BETTER ALTERNATIVE
Strong theoretical guarantees are needed.	With saliency maps, basic gradient methods can violate sensitivity and implementation invariance .	Provide stronger theoretical guarantees using integrated gradients
Spatial localisation is more important than the sensitivity at pixel-level.	With saliency maps, the gradients at pixel-level can be noisy and difficult to interpret.	Use techniques such as Grad-CAM to show important areas at a higher semantic level
Explaining black-box models	When using saliency maps, any gradient-based saliency requires differentiable neural networks .	Model agnostic explanation methods such as LIME can be used.

REAL WORLD APPLICATIONS



Medical Diagnosis

Did the model
focus on
pathology?



Autonomous Driving

Did the car stop in
front of a fire
truck?



Security Systems

What visual
evidence
triggered it?

Interpretability allows trust, debugging and interpretability.

CONCLUSION

- Saliency maps give us a way to **visually inspect how neural networks make image-based decisions**.
- **Input gradients** form mathematical basis for methods such as Vanilla Saliency, Guided Backpropagation and SmoothGrad.
- These methods help **improve transparency** and **support model debugging** in real-world scenarios.
- Despite their usefulness, saliency maps can lead to explanations that are noisy and misleading.
- They should be used critically and alongside different interpretability methods when an in-depth explanation is required.

THANK YOU FOR YOUR ATTENTION